

Rockchip Developer Guide PCIE EP Standard Card

ID: RK-KF-YF-489

Release Version: V3.2.0

Release Date: 2025-08-11

Security Level: ☐Top-Secret ☐Secret ☐Internal ☒Public

DISCLAIMER

THIS DOCUMENT IS PROVIDED "AS IS". ROCKCHIP ELECTRONICS CO., LTD.("ROCKCHIP")DOES NOT PROVIDE ANY WARRANTY OF ANY KIND, EXPRESSED, IMPLIED OR OTHERWISE, WITH RESPECT TO THE ACCURACY, RELIABILITY, COMPLETENESS, MERCHANTABILITY, FITNESS FOR ANY PARTICULAR PURPOSE OR NON-INFRINGEMENT OF ANY REPRESENTATION, INFORMATION AND CONTENT IN THIS DOCUMENT. THIS DOCUMENT IS FOR REFERENCE ONLY. THIS DOCUMENT MAY BE UPDATED OR CHANGED WITHOUT ANY NOTICE AT ANY TIME DUE TO THE UPGRADES OF THE PRODUCT OR ANY OTHER REASONS.

Trademark Statement

"Rockchip", "瑞芯微", "瑞芯" shall be Rockchip's registered trademarks and owned by Rockchip. All the other trademarks or registered trademarks mentioned in this document shall be owned by their respective owners.

All rights reserved. ©2025. Rockchip Electronics Co., Ltd.

Beyond the scope of fair use, neither any entity nor individual shall extract, copy, or distribute this document in any form in whole or in part without the written approval of Rockchip.

Rockchip Electronics Co., Ltd.

No.18 Building, A District, No.89, software Boulevard Fuzhou, Fujian, PRC

Website: www.rock-chips.com

Customer service Tel: +86-4007-700-590

Customer service Fax: +86-591-83951833

Customer service e-Mail: fae@rock-chips.com

Preface

Overview

This document provides instructions for the use of the RK PCIe standard EP development package.

Product Version

Chipset	Kernel Version
RK3588	Linux 6.1, 5.10, RT-Thread v5.1.x, v4.1.x
RK3568	Linux 6.1, 5.10, 4.19, RT-Thread v5.1.x, v4.1.x
RK1820	RT-Thread 5.1.x

Note:

- Only RT-Thread support is implemented on the RK PCIe EP side

Intended Audience

This document (this guide) is mainly intended for:

Technical support engineers

Software development engineers

Revision History

Version	Author	Date	Change Description
V1.0.0	yp.xiao, Kever Yang, Jon Lin	2023-03-07	Initial version
V1.1.0	yp.xiao	2023-05-22	Update SDK API and part of the configuration instructions
V1.2.0	yp.xiao	2023-06-15	Add EP card firmware compilation, burning, OTA upgrade instructions Add application exception instructions Update <code>struct pcie_dev_st</code> structure to support timeout configuration, PCIE speed acquisition, RC DMA configuration.
v1.3.0	Jon Lin	2023-10-07	Add pcie.bin instructions
v1.4.0	yp.xiao, Jon Lin	2023-10-23	Update SDK API Add EP acceptance criteria section Update dual storage instructions Update PCIE BIN instructions
v1.5.0	Jon Lin	2024-05-23	Update PHY Mode settings
v1.6.0	Jon Lin	2024-09-18	Add user-space test demo - <code>pcie_speed_test</code> performance reference, add support for BAR1/5
v1.7.0	yp.xiao, Jon Lin	2025-01-24	chips_connect development package adds OSAL, compatible with Windows version
v1.8.0	Jon Lin	2025-02-12	Optimize EP startup solution instructions, modify layout issues
v2.0.0	Jon Lin	2025-03-09	Add RK chip interconnection independent power supply solution instructions
v3.0.0	Jon Lin	2025-05-14	Adjust the note to EP card specifications, optimize EP Boot instructions, add data transfer principle instructions, add multi-application instructions at the user level
v3.1.0	Jon Lin	2025-07-12	Add RK1820
v3.2.0	Jon Lin	2025-08-11	Independent output standard card document

Contents

Rockchip Developer Guide PCIE EP Standard Card

1. Overview of EP Card
2. EP Card Specifications
 - 2.1 Basic Specifications
 - 2.2 Software Development
 - 2.3 Operating Systems
 - 2.4 Clock Requirements
 - 2.5 PCI Device ID
 - 2.6 Hot-Swapping
 - 2.7 Auxiliary Signals
 - 2.8 Performance Specifications
3. EP Card Operation Guide
 - 3.1 Environment Preparation
 - 3.1.1 Hardware Environment
 - 3.1.2 Software Environment
 - 3.2 Functional Verification
 - 3.2.1 Linux RC
 - 3.2.2 Windows RC
 - 3.3 Measured Performance
 - 3.4 Compilation of Samples
4. EP Card Business Basics
 - 4.1 Principle Introduction
5. Introduction to Chips Connect
 - 5.1 Principle Introduction
 - 5.2 PCIe Transfer System Architecture
 - 5.3 BUFF Management Structure and Data Transfer
 - 5.3.1 Initialization Phase
 - 5.3.2 Sending Data
 - 5.3.3 Receiving Data
 - 5.3.4 BUFF Configuration
 - 5.4 User-Level Memory Configuration
 - 5.5 Advantages of User-Level Driver
 - 5.6 Extending Multiple Independent Applications at the User Level
6. Chips Connect Configuration
 - 6.1 EP Card User-Level Device Driver
 - 6.1.1 Using-RC-DMA Mode
7. Typical Business Scenarios of Chips Connect
 - 7.1 EP Video Acceleration Card
8. Chips Connect Acceptance Criteria
 - 8.1 Development Package List
 - 8.2 Patch Update One-Click Effectiveness for "Update Package" Source Code
 - 8.3 Functional Key Log Information
 - 8.3.1 Key lspci Information on RC Side
 - 8.3.2 EP-side spl.bin Boot Log
 - 8.3.3 EP End PCIe.bin Boot Log
 - 8.4 Functionality Firmware Burning
 - 8.5 Performance User-Space Test Demo - PCIe Speed Test
 - 8.5.1 Testing Objectives
 - 8.5.2 Testing Method
 - 8.5.3 RK3588 Gen3x4 Test Results

- 8.5.4 RK3588 Gen2x1 Test Results
 - 8.5.5 RK1820 Gen2x1 Test Results
 - 8.6 Performance User Space Test Demo - PCIe Loopback Test
 - 8.6.1 Test Objective
 - 8.6.2 Testing Method
 - 8.6.3 RK3588 Gen2x1 Test Results
 - 8.6.4 RK1820 Gen2x1 Test Results
 - 8.7 Compatibility Enumeration
 - 8.8 Compatibility EP Reset and RC Re-enumeration Detection
 - 8.9 Functional OTA
- 9. Chips Connect Exception Handling
 - 9.1 RC-side x86 Platform Initialization of PCIe Reports mmap Failed, Operation Not Permitted
 - 9.2 HugePage Configuration
 - 9.2.1 X86 Platform
 - 9.2.2 RK Platform
 - 9.3 Encountering Kernel Errors or Bus Errors After Operating on the user cmd Area or ep info Area via API Interfaces
 - 9.4 Building the Windows Driver Compilation Environment
 - 9.5 Windows Application Execution Error: MSVCP140.dll and vcruntime140.dll Not Found
 - 9.6 Error Occurs When Attempting to Modify Debug Program Settings in Windows Debugging Mode. The value is protected by Secure Boot policy and cannot be modified or deleted
- 10. Chips Connect API Reference

1. Overview of EP Card

The RK PCIe EP Standard Card Development Kit is a comprehensive development package of software and hardware for RK's "PCIe EP Card" product form, based on the RK EP Card Demo Board (hereinafter referred to as the EP Card). The development kit primarily includes:

- Quick Operation Guide for EP Card
- EP Card Boot Scheme, PCIe-related hardware and software design development, and OTA upgrade content
- Basic Business of EP Card
- Business Examples under the Basic Framework of EP Card Business

2. EP Card Specifications

2.1 Basic Specifications

Main Chip	Boot Scheme	SMBus Expansion	Independent Clock Scheme
RK3568	Supports Flash Boot, Common Boot Does not support PCIe Boot	N	Supports SRNS Does not support SRIS
RK3588	Supports Flash Boot, Common Boot Does not support PCIe Boot	N	N
RK1820	Supports Flash Boot, Common Boot, EP Boot	Y	N

Explanation:

- Flash Boot includes:
 - Flash Boot - Integrated PCIe Bin and EP Boot software scheme
 - Flash Boot - Integrated PCIe Bin software scheme
- For EP boot schemes that do not support PCIe Boot, the PCIe protocol stipulates that the EP card should enter the detection mode within 20ms after the release of PERST#, followed by continuous detection, negotiation, and enumeration by the RC, which may take up to 100ms. Therefore, to enhance compatibility when the EP card interfaces with a general RC, it is recommended to choose SPI Flash as the boot storage to achieve quick startup, and SPI Nor is generally recommended as the boot storage.

2.2 Software Development

The EP card development package is essentially a PCIe product functionality demonstration development package, not a one-stop solution. Customers need to perform secondary development on PCIe and upper-layer applications based on this.

2.3 Operating Systems

EP Devices:

- EP cards are compatible with Linux, Android, and RT-Thread operating systems.

RC Devices:

- EP cards compatible with Linux operating systems for RC devices have been verified on specific RCs.
- EP cards compatible with Windows operating systems for RC devices have been verified on specific RCs, but only basic transmission functions have been implemented. Other functions require customers to develop based on the Linux driver.

2.4 Clock Requirements

It is recommended that the EP card uses a homologous reference clock scheme, that is, the EP utilizes the 100M reference clock provided by the RC end.

For EP cards that use RK chips for both RC and EP and support independent clock schemes, adopting an internal 100M clock for RC&EP can eliminate the need for an external clock generator and clock buffer chips.

2.5 PCI Device ID

The default ID for RK PCIe EP cards is the RK custom device ID. For third-party developers, the following requirements apply:

- **Internal Use Devices:** If the device is used only in a closed environment, you can use the existing VID of RK PCIe.
- **External Devices:** If the device is for sale or public use, you need to apply for a VID from PCI-SIG yourself.

2.6 Hot-Swapping

The RK PCIe Endpoint card does not support hot-swapping.

2.7 Auxiliary Signals

Auxiliary Signal	Function	Standard EP Solution	RK Chip Interconnect Independent Power Supply Solution
PRSNT#	N		
PERST#	GPIO Input	Mandatory, connected to system reset, full board reset	Mandatory, GPIO input, software reset, configure dts <code>reset-gpios;</code>
CLKREQ#	FUNCTION Input	Optional, used to extend ASPM L1SS power management, not implemented yet but recommended to reserve, if no clear demand can be left unconnected	Optional, used to extend ASPM L1SS power management, not implemented yet but recommended to reserve, if no clear demand can be left unconnected
JTAG	N		
SMBus	N		
WAKE#	GPIO Output	Mandatory, extend sleep wake-up function, not implemented yet but recommended to reserve	Mandatory, extend sleep wake-up function, configure dts <code>wake-gpios;</code>
+3.3Vaux	N		

Explanation:

- The RK chip interconnect independent power supply scheme involves scenarios where one end drops power, it is recommended to do IO protection to prevent the risk of voltage/current injection when one end drops power and the other end IO continues to output.
- Pins with GPIO function should not be configured as PCIe function in pinctrl, support any pin, no restrictions.
- Auxiliary signals with FUNCTION should select the corresponding functional pins.
- Standard EP scheme PERST# connects to system reset for full board reset requirements:
 - The PERST# signal serves as the last resort for physical reset of EP on the RC side.
 - There are special RC compatibility issues, for example: RC enumeration completes and immediately enters L2 through PERST# reset, requiring a quick reset.
 - NOTE: Since the RC's PERST is PP strong drive, and the EP's PMIC is an OD pin, regardless of whether the RC's PERST is 3.3V or 1.8V, it is necessary to drive the EP side's RESET_L through circuit conversion.

2.8 Performance Specifications

Users should clearly understand the PCIe transfer bandwidth requirements, data flow direction (RC/EP unidirectional? Bidirectional?), and whether there are fragmented control flows/large block data flows. It is recommended to:

- Read the "Measured Performance" section to understand the upper limit performance of the chips_connect development package:
- Read the "Performance | User Space Test Demo - pcie_speed_test" section to understand the transfer bandwidth for different block sizes.

3. EP Card Operation Guide

3.1 Environment Preparation

The EP card supports connection with general RC and RK RC, and this section provides a brief explanation of the basic functional verification for both scenarios.

3.1.1 Hardware Environment

RK3588 EP Board (Standard Card) Interface with RK3588 RC Reference Hardware Environment

- RC: RK3588-EVB1-V10
 - Connect the power cable, serial cable, USB-to-Type-C firmware download cable, HDMI output cable (connected to HDMI_OUT0)
- EP: RK3588-EVB4-V11
 - RC provides direct power supply, no separate power cable needed, hot-plugging is not supported, connect before powering on the RC.

RK3588 EP Board (Standard Card) Interface with General RC

- RC: PC
- EP: RK3588-EVB4-V11
 - RC provides direct power supply, no separate power cable needed, hot-plugging is not supported, connect before powering on the RC.

RK3588 Chip Interconnection Independent Power Supply Scheme Reference Hardware Environment

- RC: RK3588-EVB1-V10
 - Connect the power cable, serial cable, USB-to-Type-C firmware download cable, HDMI output cable (connected to HDMI_OUT0)
- EP: RK3588-EVB4-V12
 - EP board is independently powered, R8114 and R2005 are not populated

- It is required to connect PERST#, WAKE# signals, and use them as GPIOs
- Connect the power cable, serial cable, USB-to-Type-C firmware download cable

3.1.2 Software Environment

The loader, uboot, boot, and file system required by the chip can be obtained from the images directory of the release package, where these image files are located.

Flash Boot - Integrated PCIe Bin Software Solution

Taking the RK3588 EP board as an example, the boot files are shown in the table below:

Item	File Name	Description
RC (Not required for general RC interface)	rk3588-update-rc-evb7-v11.img	Flashed to RC Flash
EP	rk3588-update-ep-evb4-v11.img	Flashed to EP Flash

Note:

- The update.img file includes MiniLoaderAll.bin, uboot.img, boot.img, rootfs.img, and other files.

Flash Boot - Integrated PCIe Bin and EP Boot Software Solution

Taking the RK3588 EP board as an example, the boot files are shown in the table below:

Item	File Name	Description
RC (Not required for general RC interface)	rk3588-update-rc-evb7-v11.img	Flashed to RC Flash
EP	RK3588_MiniLoaderAll.bin	Flashed to EP Flash
EP	pcie_idb.img	Flashed to EP Flash
EP	normal_idb.img	Flashed to EP Flash
EP	uboot.img	Imported into EP DDR memory by RC via PCIE
EP	boot.img (including kernel+rootfs)	Imported into EP DDR memory by RC via PCIE

Note:

- For RC to download firmware to EP via PCIE, please refer to the next section on functional verification instructions.

3.2 Functional Verification

The sample code for verifying the functionality of PCIe EP is located in the examples directory of the SDK release package, including the sample codes used by EP and RC, encapsulated message communication code, and encapsulated data communication code.

3.2.1 Linux RC

The necessary bins and resources for general RC functionality verification are stored in the images directory and need to be manually copied to the RC platform (only for X86 platforms). RK RC resources are directly packaged into the firmware root directory without additional operations.

After preparing the hardware and software environment, the steps to perform a complete PCIE functionality verification are as follows. If you do not need to download firmware to the EP card, you can skip the first point directly.

1. RC downloads firmware to EP via PCIE

Step 1. The EP card's serial port baud rate is 1500000 (the RC baud rate is also 1500000). After connecting the serial port, it will continuously print the following log, indicating that the PCIE link between EP and RC is successful and waiting for RC to transmit firmware.

```
U-Boot SPL 2017.09-g614fe864bf-dirty #xyp (Apr 20 2023 - 11:19:22)
RKEP: 191277 - Link ready! Waiting RC to download Firmware:
RKEP: 191723 - Download uboot.img to BAR2+0
RKEP: 192068 - Download boot.img to BAR2+0x400000
RKEP: 192466 - Send CMD_LOADER_RUN to BAR0+0x400
RKEP: 1092850 - Waiting for FW, CMD: 0
RKEP: 2093146 - Waiting for FW, CMD: 0
RKEP: 3093443 - Waiting for FW, CMD: 0
RKEP: 4093740 - Waiting for FW, CMD: 0
RKEP: 5094036 - Waiting for FW, CMD: 0
```

Note:

If the EP card directly enters the system, it means that the flash contains a complete firmware. Please rewrite `pcie_idb.img` and `normal_idb.img` under maskrom.

Step 2. Execute `./pcie_download_test_rc 0 ./uboot.img ./boot.img` on the RC side to initiate data transfer between RC and EP. The firmware is moved to the specified memory address of the EP card via DMA, and then a startup instruction is sent to let EP boot via DDR. At this point, the EP card will print the normal boot log and enter the system.

Note:

If an exception occurs when executing the application on the RC side, please refer to `examples/pcie_download_test/README` to troubleshoot the issue.

2. Video Decoding Preview Sample

Step 1. The EP card's serial port baud rate is 1500000 (the RC baud rate is also 1500000). After connecting the serial port, executing `lspci -vvv |grep 356a` and `ls /dev/pcie-rkep*` on the RC side will show the following devices, indicating that the EP card is recognized and the RC driver is loaded normally.

```
[root@RK3588:/userdata]# lspci -vvv |grep 356a
01:00.0 Processing accelerators: Rockchip Electronics Co., Ltd Device 356a (rev 01)
[root@RK3588:/userdata]# ls /dev/pcie-rkep*
/dev/pcie-rkep0000:01:00.0
```

Step 2. After the EP card enters the system normally, the terminal will print the following log and be blocked, waiting for the RC application to start.

```
00:00:05.330 I [cmp_pcie_ep: cc_cmp_pcie_device: 390] rk pcie version: v3.0.0 - 2024-12-11
00:00:05.330 D [chipscon: cc_kernel_api_prob: 398] open rkep: /dev/pcie_ep Success.
00:00:05.330 D [chipscon: cc_kernel_api_prob: 403] probe pcie_rkep, fd = 3
00:00:05.330 I [core_pcie_ep: cc_core_pcie_dev_i: 53] bar0 size=0x4000000
00:00:05.330 I [core_pcie_ep: cc_core_pcie_dev_i: 53] bar2 size=0x4000000
00:00:05.330 I [core_pcie_ep: cc_core_pcie_dev_i: 53] bar4 size=0x1000000
00:00:05.330 E [chipscon: cc_kernel_api_reso: 141] cc_kernel_api_resource_mmap mmap failed, size=0, err=Invalid argument
00:00:05.330 E [core_pcie_ep: cc_core_pcie_dev_i: 73] no bar1
00:00:05.330 E [chipscon: cc_kernel_api_reso: 141] cc_kernel_api_resource_mmap mmap failed, size=0, err=Invalid argument
00:00:05.330 E [core_pcie_ep: cc_core_pcie_dev_i: 86] no bar5
00:00:05.330 I [pcie_dma: cc_core_pcie_dma_i: 590] - id=ffffffff
00:00:05.330 I [pcie_dma: cc_core_pcie_dma_i: 600] - id=356a1d87
00:00:05.330 I [pcie_dma: cc_core_pcie_dma_i: 609] - magic=524b4550, ver=1
.....
00:00:05.331 D [cmp_pcie_ep: pcie_task_ep_msg_t: 264] start
00:00:05.331 D [cmp_pcie_ep: stream_speed_threa: 352] stream_speed_thread tart
```

Step 3. Refer to the [User Layer Memory Configuration Section](#) to configure the memory for the RC system.

Step 4. Execute the following instructions on the RC to complete the handshake with the EP card and initiate data transfer between them.

- General RC: `./pcie_video_test_rc ./output1.h264 36 704 576 0 1`
- RK RC: `pcie_video_test_rc /root/output1.h264 36 704 576 0 1`

At this point, the RC's display device will output a resolution of 1080P60, and you can see 36 video decoding streams and the first NN human detection feature.

Note:

The RC side outputs display data through the DRM display framework. The RK RC supports it by default; just connect HDMI OUT0 to the display device. The general RC requires installing the `libdrm-dev` component and switching to another terminal without GUI by pressing `ctr+F3` before executing the test command.

Step 4. If you need to stop the Sample program, you can press the Enter key in the RC application, which will send a message to the EP card to destroy the relevant business and exit the program; the EP card will destroy the relevant business and exit the program after receiving the message. A normal business exit process is required; otherwise, the EP card business will not exit, leading to abnormal operation in the next run.

Note:

The EP side application defaults to self-start and loop start upon boot. If you need to run other test programs, you need to comment out the current test in the boot script and then restart the RC and EP. Modify /etc/init.d/S99_pcie_ep.sh and comment out the following code.

```
while true;
do
    cd /
    pcie_video_test_ep 1
done &
```

3.2.2 Windows RC


Windows RC provides drivers that only support systems below Windows 10 version 2004. Higher version systems require building a driver compilation environment to compile the drivers on their own.

Refer to the [Building Windows Driver Environment](#) section to set up the test environment, focusing only on system version and disabling driver signature enforcement steps.

After confirming the system environment, insert the EP card and power on the PC. You can see the PCI device in Device Manager, right-click on the device and select "Update Driver" -> "Browse my computer for driver software" -> set to the provided driver directory (Win-Driver-RC) -> Next to install the driver.

After installation, you can see the `Rockchip RC PCIE Driver for the rk35xx Accelerator Card` device in Device Manager.

Run the test tool with **administrator privileges**. The tool is shown in the figure below, currently supporting firmware download, rate testing, and video decoding object detection testing. If the tool reports an error, refer to the [Common Problem Handling Section](#).



After preparing the hardware and software environment, the steps to perform a complete PCIE function verification are as follows. If you do not need to download firmware to the EP card, you can skip the first point directly.

1. RC downloads firmware to EP via PCIE

Step 1. The serial port baud rate of the EP card is 1500000 (the RC baud rate is also 1500000). After connecting the serial port, it will continuously print the following log, indicating that the PCIE link between EP and RC is successful and waiting for RC to transmit firmware.

```
U-Boot SPL 2017.09-g614fe864bf-dirty #xyp (Apr 20 2023 - 11:19:22)
RKEP: 191277 - Link ready! Waiting RC to download Firmware:
RKEP: 191723 - Download uboot.img to BAR2+0
RKEP: 192068 - Download boot.img to BAR2+0x400000
RKEP: 192466 - Send CMD_LOADER_RUN to BAR0+0x400
RKEP: 1092850 - Waiting for FW, CMD: 0
RKEP: 2093146 - Waiting for FW, CMD: 0
RKEP: 3093443 - Waiting for FW, CMD: 0
RKEP: 4093740 - Waiting for FW, CMD: 0
RKEP: 5094036 - Waiting for FW, CMD: 0
```

Note:

If the EP card directly enters the system, it means that the flash contains a complete firmware. Please rewrite `pcie_idb.img` and `normal_idb.img` in the maskrom.

Step 2. Open the tool to load uboot and boot firmware, then click on firmware download. The tool will transfer the firmware to the specified memory address of the EP card via DMA and then send a start instruction to let EP boot via DDR. At this point, the EP card will print the normal boot log and enter the system.

2. Video Decoding Preview Test

Step 1: After the EP card enters the system normally, the terminal will print the following log and be blocked, waiting for the RC application to start.

```
00:00:05.330 I [cmp_pcie_ep: cc_cmp_pcie_device: 390] rk pcie version: v3.0.0 - 2024-12-11
00:00:05.330 D [chipscon: cc_kernel_api_prob: 398] open rkep: /dev/pcie_ep Success.
00:00:05.330 D [chipscon: cc_kernel_api_prob: 403] probe pcie_rkep, fd = 3
00:00:05.330 I [core_pcie_ep: cc_core_pcie_dev_i: 53] bar0 size=0x400000
00:00:05.330 I [core_pcie_ep: cc_core_pcie_dev_i: 53] bar2 size=0x4000000
00:00:05.330 I [core_pcie_ep: cc_core_pcie_dev_i: 53] bar4 size=0x100000
00:00:05.330 E [chipscon: cc_kernel_api_reso: 141] cc_kernel_api_resource_mmap mmap failed, size=0, err=Invalid argument
00:00:05.330 E [core_pcie_ep: cc_core_pcie_dev_i: 73] no bar1
00:00:05.330 E [chipscon: cc_kernel_api_reso: 141] cc_kernel_api_resource_mmap mmap failed, size=0, err=Invalid argument
00:00:05.330 E [core_pcie_ep: cc_core_pcie_dev_i: 86] no bar5
00:00:05.330 I [pcie_dma: cc_core_pcie_dma_i: 590] - id=ffffffff
00:00:05.330 I [pcie_dma: cc_core_pcie_dma_i: 600] - id=356a1d87
00:00:05.330 I [pcie_dma: cc_core_pcie_dma_i: 609] - magic=524b4550, ver=1
.....
00:00:05.331 D [cmp_pcie_ep: pcie_task_ep_msg_t: 264] start
00:00:05.331 D [cmp_pcie_ep: stream_speed_threa: 352] stream_speed_thread tart
```

Step 2: Copy the test clip to the tool directory and rename it to `output1.h264`. The current tool fixes the test clip as h264 raw stream. Set the test time, then click the video test button to run the test.

Test screen as follows:



3.3 Measured Performance

PCIe EP Transfer Performance

EP	RC	Hardware Configuration	Scenario	Maximum Transfer Rate
RK3588	RK3588	PCIe 3.0x4 Lane	EP Writes Data to RC	2.95GB/s
			EP Reads Data from RC	1.8GB/s
			EP Simultaneous Read/Write with RC	Each 1.6GB/s
RK3588	RK3588 (Enable Using RC DMA)	PCIe 3.0x4 Lane	EP Writes Data to RC	2.95GB/s
			EP Reads Data from RC	2.95GB/s
			EP Simultaneous Read/Write with RC	Each 2.85GB/s
	X86	PCIe 3.0x4 Lane	EP Writes Data to RC	2.9GB/s
			EP Reads Data from RC	2.5GB/s
			EP Simultaneous Read/Write with RC	Each 2.4GB/s
	Windows	PCIe 3.0x4 Lane	EP Writes Data to RC	650MB/s
			EP Reads Data from RC	700MB/s
			EP Simultaneous Read/Write with RC	Each 320MB/s

Performance will continue to be optimized. Windows does not currently support chain DMA testing, and the transfer rate is significantly affected by the memory structure; higher rates are achieved with large blocks of contiguous memory.

3.4 Compilation of Samples

For compilation and usage instructions of each sample, please refer to `examples/*/README`.

4. EP Card Business Basics

4.1 Principle Introduction

The EP card business primarily involves transferring source data and target data processed by the processor through the PCIe link between upstream and downstream. The main resources include:

RC Side

- Manipulate PCIe EP devices through the open-source standard library libpci.
- Complete with the pcie-rkep driver:
 - BAR mapping, recommended as a control flow transfer method
 - DMA transfer, recommended as a data block transfer method
 - Handle MSI interrupts, with the user layer actively polling to wait for an increase in response rate, but involving multiple handshakes, real-time performance is average
 - Trigger ELBI interrupts
 - Apply for DMA memory
 - Support for multi-process, such as supporting the operation of multiple pcie_speed_test routines

EP Side

- Complete with the pcie_ep driver:
 - Default class id 0x1200 (Processing accelerators)
 - BAR inbound physical address space mapping
 - DMA transfer
 - Send MSI interrupts
 - Handle ELBI interrupts, with the user layer actively polling to wait for an increase in response rate, but involving multiple handshakes, real-time performance is average
 - Support for multi-threading
- Hardware resources such as VDEC, VENC, RGA, GPU, NPU can be called at the user layer through Rockit or RKNN.

5. Introduction to Chips Connect

5.1 Principle Introduction

Chips Connect development kit is a user-level development kit built based on the EP card business foundation, primarily providing interface references for commonly used functions such as RC, EP uplink and downlink data transmission, and message notifications, while also constructing typical business scenarios for reference.

BAR Space Allocation (common platforms)

For general schemes with sufficient memory allocation, BAR space allocation during SPL and Kernel stages:

BAR Space	Start Address	End Address	Function
BAR0	0	4KB	EP reserved space, driver version information, interrupt status, etc.
	12KB	16KB	DEBUG function-related configuration information.
	16KB	128KB	EP device status information.
	128KB	1M	User-defined.
	1M	4M	Framework command interaction and buff management.
BAR1 (Default Disabled)	0	1MB	User-defined, how to enable refer to the "BAR Size Configuration" section.
BAR2 (BAR2/3 Bar Repair)	0	64MB	Using RC DMA.
BAR4	0	4MB	Wired register area, used for operating DMA.
BAR5 (Default Disabled)	0	1MB	User-defined, how to enable refer to the "BAR Size Configuration" section.

Explanation:

- EP end BAR user-defined space can be used for secondary development expansion, the rest of the BAR space is the space required for the business code, for example:
 - BAR1 is disabled by default, all space is user-defined, and when enabled, it is used for control flow.
 - BAR5 is disabled by default, all space is user-defined, and when enabled, it is used for control flow.
- The EP end BAR inbound CPU address is fixed by the "BAR Mapping Configuration" by default and supports dynamic adjustment on the RC end.

BAR Space Allocation (memory-constrained platforms)

For special schemes with tight memory allocation, BAR space allocation during SPL and Kernel stages:


BAR Space	Start Address	End Address	Function
BAR0	0	Reserved	Reserved space, default is 0.
	R	R+2048B	EP reserved space, driver version information, interrupt status, etc.
	R+2048B	R+2560B	DEBUG function-related configuration information.
	R+2056B	R+3072B	EP device status information.
	R+3072B	R+4KB	User-defined.
	R+4KB	R+64KB	Framework command interaction and buff management.
BAR1 (Default Disabled)		1MB	User-defined, how to enable refer to the "BAR Size Configuration" section.
BAR2 (BAR2/3 Bar Repair)	0	64MB	Using RC DMA.
BAR4	0	4MB	Wired register area, used for operating DMA.
BAR5 (Default Disabled)	0	1MB	User-defined, how to enable refer to the "BAR Size Configuration" section.

Explanation:

- Reserved space: The minimum alignment of BAR is 1MB, and it is usually required to apply for 1MB of memory for mapping. However, in environments with tight memory, reserved processing is allowed, and the chips connect business uses the reserved memory.
- EP end BAR user-defined space can be used for secondary development expansion, the rest of the BAR space is the space required for the business code, for example:
 - BAR1 is disabled by default, all space is user-defined, and when enabled, it is used for control flow.
 - BAR5 is disabled by default, all space is user-defined, and when enabled, it is used for control flow.
- The EP end BAR inbound CPU address is fixed by the "BAR Mapping Configuration" by default and supports dynamic adjustment on the RC end.


5.2 PCIe Transfer System Architecture

PCIe RC and EP System Block Diagram



5.3 BUFF Management Structure and Data Transfer

BUFF Management



Here is a brief explanation from the RC perspective on the initialization to a complete read-write process:

5.3.1 Initialization Phase

Primarily, two tasks are performed:

1. Initialize the corresponding PCIe device based on `VENDOR_ID`, `DEVICE_ID`, and bus address.
Internally, this is mainly achieved by iterating through the system's PCIe devices via the `libpci` interface and mapping the BAR space to the application layer.
It supports calling the `rk_pcie_device_init` interface repeatedly with different device IDs to initialize the device.
2. Call `rk_pcie_task_create` to create a task for initializing memory and BUFF status. The number and size of BUFFs can be set through the `struct pcie_task_attr_st` structure. The configuration information is passed to the EP device through the BAR0 space and waits for the EP device to complete

initialization. The EP retrieves the task creation message through `rk_pcie_get_task_msg` and then calls task initialization, facilitating an initialization configuration interaction between both sides.

The RC-side application, based on compilation configuration information, selects Hugepage/RKEP/RKDRM methods to allocate memory. Each EP card device will generate a corresponding `/dev/pcie-rkep*` node on the RC side and allocate separate contiguous large memory (requires corresponding configuration to be enabled).

After receiving the configuration information, the EP initializes the BUFF status and calls the interface to apply for physically contiguous memory, supporting Hugepage/RKDRM methods for memory allocation.

The BUFF information for both RC and EP is ultimately assigned to a specific area of BAR0, which is used for BUFF status management.

```
#define VENDOR_ID            0x1d87
#define DEVICE_ID            0x356a

struct pcie_dev_bus dev_bus;
int ret;
//Retrieve the system-matched PCI device bus address
ret = rk_pcie_get_pci_bus_addresses(VENDOR_ID, DEVICE_ID, &rc_ctx.dev_bus);
if (ret != 0) {
    printf("get pci bus address fail\n");
    return -1;
}

printf("ep dev max num : %d\n", rc_ctx.dev_bus.dev_max_num);
if (rc_ctx.dev_bus.dev_max_num == 0) {
    printf("no ep device, exit ... \n");
} else {
    rc_ctx.dev_attr.using_rc_dma = 0;
    rc_ctx.dev_attr.rc_dma_chn = 0;
    rc_ctx.dev_attr.enable_speed = 1;
    rc_ctx.dev_attr.vendor_id = VENDOR_ID;
    rc_ctx.dev_attr.device_id = DEVICE_ID;
    rc_ctx.dev_attr.use_dma_only_in_rc = 0;
    rc_ctx.dev_attr.mem_type = E_MEM_TYPE_HUGEPAGE;
    memcpy(rc_ctx.dev_attr.pci_bus_address, pci_bus_address, strlen(pci_bus_address));

    //Select the first matching device for initialization
    rc_ctx.rk_handle = rk_pcie_device_init(&rc_ctx.dev_attr);
    //Retrieve EP device mode
    rk_pcie_get_ep_mode(rc_ctx.rk_handle, &rc_ctx.devmode);
    printf("ep device init, dev mode: %d, submode: %d\n", rc_ctx.devmode.mode,
rc_ctx.devmode.submode);
}
```

5.3.2 Sending Data

Sending data can be divided into three steps:

1. Check if there is an available BUFFER.
2. Obtain the available BUFFER and write data into it.
3. Send the BUFFER to the other end.


Taking the sending of a BUFFER as an example, initializing the configuration `send_buff_num = 5` indicates that there can be at most 5 sending BUFFERS in the system. The sender calls the interface to obtain an available BUFFER, and can then access the physical address and virtual address of the BUFFER, as well as the size of the BUFFER, through the properties of `struct pcie_buff_node`. After that, business operations can be performed. Finally, the sending interface is called to transmit the BUFFER data to the other end. Inside the sending interface, it will first flush the cache of the BUFFER to ensure that the physical memory contains the most recent data, and then call the DMA interface to transfer the data to the other end. The size of the data transferred by DMA is `node->size`, which should not exceed the size of the BUFFER.

Note: The BUFFER sent to the other end needs to be released by the other end before it can become available again. If the other end does not release it, the entire process will be blocked.


`rk_pcie_get_buff` does not support blocking mode temporarily, so it is necessary to call sleep to wait in a loop.

```
struct pcie_buff_node *node = rk_pcie_get_buff(rk_handle, E_SEND);
if(node) {
    node->size = rk_pcie_get_buff_max_size(rk_handle, E_SEND);
    //operate on buff, write data
    rk_pcie_send_buff(ctx->rk_handle, node);
} else {
    usleep(1000);
}
```

RC Sending Data to EP Principle:



EP Sending Data to RC Principle:



5.3.3 Receiving Data


Receiving data can be divided into three steps:

1. Check if there is an available BUFFER.
2. Acquire the available BUFFER and read the data.
3. Release the BUFFER.


Receiving BUFFER and sending BUFFER have similar internal implementations. The number of BUFFERS in the system is configured based on the initialization of `recv_buff_num`. Specific behavior can be referenced in the section on sending data. The size of the available data in the BUFFER is `node->size`.

```
struct pcie_buff_node *node = rk_pcie_get_buff(rk_handle, E_RECV);
if(node) {
    //operate on buff to read data
    rk_pcie_release_buff(rk_handle, node);
} else {
    usleep(1000);
}
```

EP Data Reception Principle:



RC Data Reception Principle:



5.3.4 BUFF Configuration

1. BUFF Configuration for RC to EP Transmission

In the decoding card scenario, the RC end sends raw streams to the EP, with relatively low data volume per channel. A BUFF size of `4KB*video_channel` can be allocated. Each video stream transmits a size of 4KB per transfer, ensuring that each transmission can update the data and that the data volume is not too small to waste the PCIe bandwidth.

It is recommended to configure 3-5 BUFFs for rotation use. At least three BUFFs must be configured; if additional processing is required for the BUFF, the number of BUFFs should be increased, otherwise, it will cause rotation to stall and affect the transmission rate.

Multi-threaded transmission can appropriately increase the number of BUFFs.

2. BUFF Configuration for EP to RC Transmission.

The data sent from EP to RC is usually decoded data, which is generally larger, such as a single frame of 1080p decoded into YUV422 data being approximately 3MB in size. A single BUFF can be configured to be 3MB-6MB in size, transmitting 1 or 2 channels of decoded data per transfer. Since multiple BUFFs are needed for rotation, a single BUFF should not be configured too large to avoid occupying too much memory space.

5.4 User-Level Memory Configuration

Both RC and EP support various memory allocation methods:

- RC side:

1. Memory allocation via the pcie-rkep driver, SDK obtains it through mmap mapping.

Enable configuration: `CONFIG_PCIE_FUNC_RKEP_USERPAGES=y` (disabled by default)

Modify memory size: Modify the macro `RKEP_USER_MEM_SIZE` in the driver `pcie-rkep.c` (default is 64M)

2. Memory allocation through HugePages. (Recommended method)

3. Memory allocation via the DRM driver, requiring pre-allocation of CMA in the dts.

Modify DMA memory size in device tree: `arch/arm64/boot/dts/rockchip/rk3588-linux.dtsi`

```
@@ -71,7 +71,7 @@ reserved-memory {
cma {
    compatible = "shared-dma-pool";
    reusable;
    reg = <0x0 0x50000000 0x0 0x10000000>; // Starting from address 0x50000000, size
is 256MB
    linux,cma-default;
};
```

4. Dummy method, where the data buffer is managed by the user, and the framework only manages the rotation node without allocating the data buffer. Currently, only supports Windows systems.

- EP side:

1. Memory allocation via the DRM driver, requiring pre-allocation of CMA in the dts.

Modify DMA memory size in device tree: `arch/arm64/boot/dts/rockchip/rk3588-linux.dtsi`

```
@@ -71,7 +71,7 @@ reserved-memory {
    cma {
        compatible = "shared-dma-pool";
        reusable;
        reg = <0x0 0x50000000 0x0 0x10000000>; // Starting from address 0x50000000, size
        is 256MB
        linux,cma-default;
    };
```

2. Memory allocation through HugePages. (Recommended method)

Different memory allocation methods can be specified through attributes during Task initialization.

- RC side: `E_MEM_TYPE_HUGEPAGE` / `E_MEM_TYPE_RKDRM` / `E_MEM_TYPE_RKEP`
- EP side: `E_MEM_TYPE_HUGEPAGE` / `E_MEM_TYPE_RKDRM`

Note:

Allocation using HugePages requires system configuration, which can be configured by referring to the [HugePage section](#).

5.5 Advantages of User-Level Driver

The primary consideration for elevating the PCIe data transfer service driver to the user level is to facilitate usage and optimize performance.

Usage aspects:

- Based on the PCIe device driver, it facilitates portability across different platforms;
- Facilitates direct integration with applications, allowing direct calls to RK's provided rockit interface or third-party gstreamer interfaces, etc.;

In terms of performance optimization compared to kernel interface drivers, the application layer driver mainly offers the following advantages:

- Replaces interrupts with polling to reduce the CPU mode switching overhead caused by frequent interrupts;
- Reduces the CPU mode switching overhead caused by kernel and userspace transitions;
- It is recommended to bind to a specific CPU to reduce multi-core scheduling overhead;
- Eliminates the memory copy overhead between kernel space and userspace, achieving memory zero-copy;
- Utilizes system Huge Page memory to reduce TLB cache miss.

5.6 Extending Multiple Independent Applications at the User Level

RK chips_connect user-level application foundation is based on Task:


- Tasks are independent of each other
 - RC user messages are independent, RC Task supports multi-process development
 - EP user messages are managed uniformly, EP Task does not support multi-process
- Multi-threading is supported within a Task
- Task ID is assigned on a first-come, first-served basis and cannot specify a specific ID

Since the default reference case application is not bound to Task and other resources, it is not possible to implement multiple independent applications in parallel. It is recommended to extend multiple independent applications based on user messages with application IDs, and Task ownership. For example, RK RT-Thread chips_connect multi-application solution:

1. Define user message rules, add app management:

```
struct task_user_cmd_st {  
    unsigned int app_id;  
    unsigned int cmd;  
    unsigned char data[4];  
};
```

2. EP application registers the application, unique app_id
3. RC sends messages, unique app_id
4. EP main program triggers application message callbacks based on app_id
5. EP application processes messages, including applying for/releasing Task



6. Chips Connect Configuration

6.1 EP Card User-Level Device Driver

The EP Card User-Level Device Driver refers to the EP function driver running on the RC side, including kernel modules and user-level APPs.

6.1.1 Using-RC-DMA Mode

The Using-RC-DMA mode operates exclusively when the EP card is connected to RK RC devices, allowing the PCIe bandwidth to be maximized.

By default, only the EP-side DMA is utilized. The principle of the Using-RC-DMA mode is to leverage the DMA on the RC side as well. When the EP writes data to the RC, it calls upon the EP's DMA; when the RC writes data to the EP, it calls upon the RC's DMA, without interference.

The RC-side DMA is limited to transferring data to physical spaces that have memory inbound mapping to BAR. Therefore, BAR2 is allocated for RC DMA usage. The default size is 64M, which can be modified in the EP-side dts configuration. This memory is used only when the RC writes data to the EP, specifically during the initialization phase when the `send_buff_size` is configured, and it must not be smaller than `send_buff_size`.

Enabling the Mode

No modifications are required for the RC and EP device drivers; only the parameter configuration during the initialization phase of the RC upper-layer application is needed.

```
dev_attr.using_rc_dma = 1; //Enable RC DMA
dev_attr.rc_dma_chn = 1; //Channel 1, supports 0/1
```

After successful enabling, the following LOG can be observed in the RC application:

```
rk pcie version: v2.0.0 - 202300927
using rc dma: 1, chn: 1 //Enable RC DMA, Channel 1
open rkrmd: /dev/pcie-rkep-0000:01:00.0 Success.
```

When the EP application creates a Task, the following LOG will be displayed:

```
task create id: 1
using rc dma: 1
```


Recommended Configuration

- Single EP scenario: It is recommended to enable `using_rc_dma` and configure it to Channel 0.
- Multiple EP scenarios: It is recommended to disable `using_rc_dma`, using the EP's own DMA to transfer data, which can maintain a roughly equal transmission rate for each EP.

7. Typical Business Scenarios of Chips Connect

7.1 EP Video Acceleration Card

The video acceleration card transfers RC's video stream to EP via the PCIE bus for decoding and processing, and then transmits it back to the main chip for display via the PCIE bus. The basic data flow processing is shown in the figure below:



The stream sender first obtains the stream data from the video stream, copies it into the prepared send stream buffer, and then sends the stream data to the rcv stream buffer on the PCIE opposite end through PCIE DMA. The opposite end then extracts the stream and processes it through VDEC into VPSS, which outputs two streams of data. One stream is sent to the GPU for synthesis, and the other stream is sent to RKNN for processing. The synthesized data is copied into the prepared send stream buffer, and then the data is sent to the rcv stream buffer on the PCIE opposite end through PCIE DMA. The opposite end extracts the data and sends it to the VO module for display.

The SDK has implemented the stream buffers for both the sender and receiver, which can be used by configuring parameters during the initialization phase. Users only need to focus on the business process.

8. Chips Connect Acceptance Criteria

It is recommended to follow the process below to confirm that the software, hardware, and basic business processes are functioning correctly when using the PCIe EP development kit for the first time, before starting custom feature development.

8.1 Development Package List

Directory Structure Introduction

```
.
├─ bootloader           # PCIe EP bootloader source code project
├─ build                # build compilation scripts
├─ core
```

```

|   ├── api                # SDK external API, including RC and EP
|   └── component          # PCIe components, mainly responsible for buf management,
message notification, etc.
|   └── framework          # Framework layer, mainly responsible for interacting with
OS, providing system capabilities and PCIe functions
└── docs                  # PCIe EP SDK documentation introduction
└── examples
|   ├── pcie_camera_test   # EP card typical business, capturing EP Camera data
processing, only for Android
|   ├── pcie_speed_test    # EP card typical business, PCIe transfer rate test Demo
|   ├── pcie_video_test    # EP card typical business, video card Demo
|   ├── pcie_download_test # EP card RC download firmware startup solution Demo
|   ├── pcie_auto_test     # EP card functional automated testing Demo
|   └── windows            # EP card Windows testing Demo
└── images                # Functional verification firmware
└── patch*                # Patch directory
└── third-party           # Third-party dependencies

```

Introduction to the main directories of the development package:

The core/framework directory implements basic functions, including memory allocation management, logging functions, linked_list, DMA, PCIe CONFIG, and other RC and EP common components.

- framework/core/pcie: Implements DMA transfer, responsible for interacting with chip DMA configuration.
- framework/core/mem: Implements memory allocation and release under different configurations.
- framework/osal: Encapsulates the implementation of OS interfaces, shielding the differences between different OS from the upper layer.
- framework/base: Implements basic data structures, such as linked lists, queues, etc.

The core/component directory encapsulates the implementation of RC and EP interfaces, internally implementing buf management, message notification, data transfer, etc., and externally providing specific APIs for development without calling through the OS. This directory consists of three parts:

- ep_pcie: Responsible for buf management, message notification, and data transfer related implementations for the PCIe EP part.
- rc_pcie: Responsible for buf management, message notification, and data transfer related implementations for the PCIe RC part.
- cc_pcie_type.h and cc_cmp_pcie_type.h: Provide unified data structures and data types for external communication between RC and EP.

The core/api directory encapsulates the implementation of different system user interface layers, providing specific APIs for user development. This directory consists of four parts:

- ep: RK EP only needs to support the arm environment, so this part only encapsulates the component/ep_pcie interface without implementing additional logic.
- rc/windows: The Windows system is developed using driver forms, providing ioctl calls for users.
- rc/other: Other systems besides Windows do not need to use driver development, so this part only encapsulates the component/rc_pcie interface without implementing additional logic.

build: The directory implements the compilation and construction of the SDK, supporting Linux, Android, and Windows.

docs: Standard EP board documentation.

examples: Provides different test samples, sample introductions, and operation instructions. Please refer to the README in the corresponding directory.

images: The directory contains firmware and samples that can be directly used for basic functional testing.

patch*: Patches corresponding to the SDK package, supporting Linux and Android systems. Patch instructions can be found in `patch*/README`.

8.2 Patch Update One-Click Effectiveness for "Update Package" Source Code

Testing Objective

Standard card driver version upgrade, supporting one-click addition of patches in the "Update Package" directory.

Testing Method

RK3588:

```
./ep_patch.sh release patch-kernel5.10/Update Package ./path-to-dst/
```

RK3568:

```
./ep_patch.sh release patch-kernel4.19/Update Package ./path-to-dst/
```

Additional Notes:

- Generating patch packages:

```
./ep_patch.sh release ../../ patch-kernel5.10/Update Package
cd ../../kernel/
git checkout develop-4.19
cd -
./ep_patch.sh kernel ../../ patch-kernel4.19/Update Package
```

8.3 Functional Key Log Information

8.3.1 Key lspci Information on RC Side

Confirm the following information:

- bar is enabled and properly allocated

- MSI is enabled and properly allocated
- LnkCap, Gen3x4 support, ASPM not supported
- LnkSta, Gen3x4 status
- Kernel driver in use: pcie-rkep

Log example:

```
01:00.0 Class 1200: Device 1d87:356a (rev 01)
    Control: I/O- Mem+ BusMaster+ SpecCycle- MemWINV- VGASnoop- ParErr- Stepping- SERR-
FastB2B- DisINTx+
    Status: Cap+ 66MHz- UDF- FastB2B- ParErr- DEVSEL=fast >TAbort- <TAbort- <MAbort-
>SERR- <PERR- INTx-
    Latency: 0
    Interrupt: pin A routed to IRQ 157
    Region 0: Memory at f0400000 (32-bit, non-prefetchable) [size=4M]
    Region 2: Memory at 900000000 (64-bit, prefetchable) [size=64M]
    Region 4: Memory at f0200000 (32-bit, non-prefetchable) [size=1M]
    Capabilities: [40] Power Management version 3
        Flags: PMEClk- DSI- D1+ D2+ AuxCurrent=375mA PME(D0+,D1+,D2-,D3hot+,D3cold-)
        Status: D0 NoSoftRst+ PME-Enable- DSel=0 DScale=0 PME-
    Capabilities: [50] MSI: Enable+ Count=4/32 Maskable+ 64bit+
        Address: 00000000fe670040 Data: 0000
        Masking: ffffffff0 Pending: 00000000
    Capabilities: [70] Express (v2) Endpoint, MSI 08
        DevCap: MaxPayload 256 bytes, PhantFunc 0, Latency L0s unlimited, L1
unlimited
        ExtTag+ AttnBtn- AttnInd- PwrInd- RBE+ FLReset- SlotPowerLimit
0.000W
        DevCtl: Report errors: Correctable- Non-Fatal- Fatal- Unsupported-
        RlxdOrd+ ExtTag+ PhantFunc- AuxPwr- NoSnoop-
        MaxPayload 256 bytes, MaxReadReq 256 bytes
        DevSta: CorrErr- UncorrErr- FatalErr- UnsuppReq- AuxPwr- TransPend-
        LnkCap: Port #0, Speed 8GT/s, Width x4, ASPM not supported, Exit Latency L0s
<4us, L1 <16us
        ClockPM- Surprise- LLActRep- BwNot- ASPMOptComp+
        LnkCtl: ASPM Disabled; RCB 64 bytes Disabled- CommClk-
        ExtSynch- ClockPM- AutWidDis- BWInt- AutBWInt-
        LnkSta: Speed 8GT/s, Width x4, TrErr- Train- SlotClk+ DLActive- BWMgmt-
ABWMgmt-
        DevCap2: Completion Timeout: Not Supported, TimeoutDis+, LTR+, OBFF Via
message/WAKE#
        DevCtl2: Completion Timeout: 50us to 50ms, TimeoutDis-, LTR+, OBFF Disabled
        LnkCtl2: Target Link Speed: 8GT/s, EnterCompliance- SpeedDis-
        Transmit Margin: Normal Operating Range, EnterModifiedCompliance-
ComplianceSOS-
        Compliance De-emphasis: -6dB
        LnkSta2: Current De-emphasis Level: -6dB, EqualizationComplete+,
EqualizationPhase1+
        EqualizationPhase2+, EqualizationPhase3+, LinkEqualizationRequest-
        Capabilities: [b0] MSI-X: Enable- Count=128 Masked-
        Vector table: BAR=4 offset=00020000
        PBA: BAR=4 offset=00028000
```

```

Capabilities: [100 v2] Advanced Error Reporting
    UESSta: DLP- SDES- TLP- FCP- CmpltTO- CmpltAbrt- UnxCmplt- RxOF- MalfTLP-
ECRC- UnsupReq- ACSViol-
    UEMsk: DLP- SDES- TLP- FCP- CmpltTO- CmpltAbrt- UnxCmplt- RxOF- MalfTLP-
ECRC- UnsupReq- ACSViol-
    UESvrt: DLP+ SDES+ TLP- FCP+ CmpltTO- CmpltAbrt- UnxCmplt- RxOF+ MalfTLP+
ECRC- UnsupReq- ACSViol-
    CESTa: RxErr- BadTLP- BadDLLP- Rollover- Timeout- NonFatalErr-
    CEMsk: RxErr- BadTLP- BadDLLP- Rollover- Timeout- NonFatalErr+
    AERCap: First Error Pointer: 00, GenCap+ CGenEn- ChkCap+ ChkEn-
Capabilities: [148 v1] #19
Capabilities: [168 v1] Address Translation Service (ATS)
    ATSCap: Invalidate Queue Depth: 00
    ATSCtl: Enable-, Smallest Translation Unit: 00
Capabilities: [178 v1] Page Request Interface (PRI)
    PRICtl: Enable- Reset-
    PRISa: RF- UPRGI- Stopped+
    Page Request Capacity: 00000001, Page Request Allocation: 00000000
Capabilities: [188 v1] Latency Tolerance Reporting
    Max snoop latency: 0ns
    Max no snoop latency: 0ns
Capabilities: [190 v1] L1 PM Substates
    L1SubCap: PCI-PM_L1.2+ PCI-PM_L1.1+ ASPM_L1.2+ ASPM_L1.1+ L1_PM_Substates+
    PortCommonModeRestoreTime=10us PortTPowerOnTime=10us
Capabilities: [1a0 v1] #16
Capabilities: [1d0 v1] Vendor Specific Information: ID=0002 Rev=4 Len=100 <?>
Capabilities: [2d0 v1] Vendor Specific Information: ID=0006 Rev=0 Len=018 <?>
Capabilities: [2e8 v1] #15
Kernel driver in use: pcie-rkep

```

8.3.2 EP-side spl.bin Boot Log

Verify the following information:

- Kernel phase prints "already linkup"

Log example:

```

... # SPL key print

U-Boot SPL board init
RKEP: Init PCIe EP
RKEP: 188259 - PHY Mode 0x4
RKEP: 189537 - RKEP: GRF:904=7, a04=7...
RKEP: 190138 - RKEP: GRF:904=7, a04=f...
RKEP: 190487 - RKEP: GRF:904=f, a04=f...
RKEP: 190827 - RefClock in common clock_mode
RKEP: 200593 - BAR0: 0x3c000000
RKEP: 200864 - BAR2: 0x010000000
RKEP: 201247 - init PCIe fast Link up
RKEP: 218243 - Link up 230011
RKEP: 259699 - Done

```

```
... # Kernel key print
```

```
[00-06-34.831][ 3.148198] rk-pcie-ep fe150000.pcie: bar0: assigned [0x3c000000-3c3ffffff]
[00-06-34.834][ 3.149015] rk-pcie-ep fe150000.pcie: bar2: assigned [0x40000000-43ffffff]
[00-06-34.834][ 3.149027] rk-pcie-ep fe150000.pcie: no vpcie3v3 regulator found
[00-06-34.837][ 3.149046] rk-pcie-ep fe150000.pcie: already linkup
[00-06-34.837][ 3.149320] rk-pcie-ep fe150000.pcie: register misc device pcie_ep
```

8.3.3 EP End PCIE.bin Boot Log

Confirm the following information:

- Record the running time of pcie.bin

Version	Duration
V1.00	360ms

- SPL phase prints "already link up"
- Kernel phase prints "already linkup"

Log example:

```
# Key prints of pcie.bin
PCIE_EP Release Time: Jan  2 2025 V2.20
Gen3x4
phyMode4
HW
controlelr initialed
GRF:904=0007, a04=0007...
GRF:904=004F, a04=004F...
phy locked
phy initialed
vid=0x1D87 did=0x356A
ms=000C
Linking, ltssm:
00000001
00000002
00000007
00210022
00210023
00230011
Link up
rst_fun1
Link stable, ltssm
00230011
out

... # Key prints of SPL
```

```
U-Boot SPL board init
RKEP: Init PCIe EP
RKEP: already link up

... # Key prints of kernel

[00-06-34.831][ 3.148198] rk-pcie-ep fe150000.pcie: bar0: assigned [0x3c000000-3c3ffffff]
[00-06-34.834][ 3.149015] rk-pcie-ep fe150000.pcie: bar2: assigned [0x40000000-43ffffff]
[00-06-34.834][ 3.149027] rk-pcie-ep fe150000.pcie: no vpcie3v3 regulator found
[00-06-34.837][ 3.149046] rk-pcie-ep fe150000.pcie: already linkup
[00-06-34.837][ 3.149320] rk-pcie-ep fe150000.pcie: register misc device pcie_ep
```

8.4 Functionality Firmware Burning

Please refer to the "Single EP Card Interface Universal RC" section to confirm whether it is a product form that requires a quick start, and confirm the storage solution used and perform the corresponding PC upgrade burning verification, mainly including:

- Single Storage
- Dual Storage
- EP Boot

For detailed information, refer to the "EP Card Firmware Compilation" section.

8.5 Performance User-Space Test Demo - PCIe Speed Test

8.5.1 Testing Objectives

Verify the business fundamentals, which mainly include:

- Wired DMA to EP transfer on the RC side, DMA transfer rate, and bidirectional DMA transfer
- Local DMA to RC transfer on the RC side, DMA transfer rate, and bidirectional DMA transfer
- BAR mapping, normal access
- ELBI operates normally, supporting RC to trigger interrupt event notifications to EP
- MSI operates normally, supporting EP to trigger interrupt event notifications to RC

Based on requirements, verify:

- If a DMA-supported RK RC is selected, it supports Local DMA to EP transfer on the RC side, and DMA transfer rate
- Firmware download

8.5.2 Testing Method

For detailed reference, please consult the document [examples/pcie_speed_test/README.md](#).

8.5.3 RK3588 Gen3x4 Test Results

Test Equipment

RC: RK3588-EVB1-LP4X-V10 PCIe 3.0 4 lanes with main core at 2.4G and auxiliary core at 1.8G, CPU/Dram set to high frequency

EP: RK3588-EVB4-LP4X-V10 PCIe 3.0 4 lanes with main core at 2.4G and auxiliary core at 1.8G, CPU/Dram set to high frequency

Standard EP Version

Rockchip_PCIE_EP_Standard_Card_20231215.tar.gz

Test Parameters:

- Utilizing EP DMA engine
- task_num = 1

Test Results

DMA Size	64B	256B	4K	64KB	1MB	4MB
RC Read (DMA Write)	0.16MB/S	0.65MB/S	10.4MB/S	165.8MB/S	1898.7MB/S	3271.3MB/s
RC Write (DMA Read)	0.17MB/S	0.73MB/S	11.0MB/S	171.1MB/S	1630.6MB/S	1735.6MB/s

Explanation:

- If the test is initiated by RC and uses the default EP DMA engine, the write test is effectively an EP DMA read, and the read test is effectively an EP DMA write.
- The smaller the test granularity, the more scheduling is involved, leading to poorer performance, but overall performance can be improved by using multi-threading.

8.5.4 RK3588 Gen2x1 Test Results

Test Equipment

RC: RK3588-EVB1-LP4X-V10 PCIe 3.0 4 lanes with 2.4G for the main core and 1.8G for the small core.

EP: RK3588-EVB4-LP4X-V10 PCIe 3.0 4 lanes (reduced to Gen2 1 lane) with 2.4G for the main core and 1.8G for the small core.

Standard EP Version

Rockchip_PCIE_EP_Standard_Card_20231215.tar.gz

Test Parameters:

- Utilizing EP DMA engine
- task_num = 1

Test Results

DMA Size	64B	256B	4K	64KB	1MB	4MB
Write	0.16MB/S	0.64MB/S	10.2MB/S	147.6MB/S	394.0MB/S	396.2MB/s
Read	0.17MB/S	0.70MB/S	11.0MB/S	149.2MB/S	384.0MB/S	384.1MB/s

Explanation:

- If the test is initiated by RC and uses the default EP DMA engine, the write test actually corresponds to EP DMA read, and the read test actually corresponds to EP DMA write.
- The smaller the test granularity, the more scheduling is involved, leading to poorer performance, but overall performance can be improved by utilizing multi-threading.

8.5.5 RK1820 Gen2x1 Test Results

Test Equipment

RC: RK3588-EVB1-LP4X-V10 PCIe 3.0 4 lanes, big core 2.4G, little core 1.8G

EP: RK1820 PCIe 2.0 1 lane

Standard EP Version

RC - Rockchip_PCIE_EP_Standard_Card_20250523_V3.1.0.tar.gz

EP - RK RT-Thread SDK develop-v5.1.x branch

Test Parameters:

- Using EP DMA engine
- task_num = 1

RK1820 DRAM Memory Test Results

DMA Size	64B	256B	4K	8K	64KB	512KB
RC Write (EP DMA Read)	3MB/S	10MB/S	136MB/S	188MB/s	303MB/s	317MB/s
RC Read (EP DMA Write)	3MB/S	11MB/S	143MB/S	237MB/S	363MB/s	382MB/s

Explanation:


- If the test is initiated by RC and uses the default EP DMA engine, the write test is actually EP DMA read, and the read test is actually EP DMA write.
- The smaller the test granularity, the more scheduling is involved, the worse the performance, but overall performance can be improved by opening multiple threads.
- Compared to EP Linux OS, RT-Thread OS is a real-time system, allowing user-level programs to quickly

respond to interrupts, so small granularity data transfer still maintains a high throughput.

8.6 Performance User Space Test Demo - PCIe Loopback Test

8.6.1 Test Objective

Count the typical packet sending and receiving scenarios, and measure the response time duration of the serial process.



8.6.2 Testing Method

For detailed reference, please consult the document located at [examples/pcie_speed_test/README.md](#).

8.6.3 RK3588 Gen2x1 Test Results

Test Equipment

RC: RK3588-EVB1-LP4X-V10 PCIe 3.0 4 lanes with main core at 2.4G and secondary core at 1.8G

EP: RK3588-EVB4-LP4X-V10 PCIe 3.0 4 lanes (reduced to Gen2 1 lane) with main core at 2.4G and secondary core at 1.8G

Standard EP Version

Rockchip_PCIE_EP_Standard_Card_20231215.tar.gz

Test Parameters:

- Utilizing EP DMA engine
- task_num = 1

Test Results

DMA Size	4K
min/max	105us/724us

Explanation:

- If the test is initiated by RC and the default EP DMA engine is used, the write test is effectively an EP DMA read, and the read test is effectively an EP DMA write.
- The test results are statistical values after 1 second of testing, excluding the bias effects of the initialization process.
- The test results represent the best case scenario under continuous packet transmission; adjust intervals and other loads as needed to construct a worst case scenario.

8.6.4 RK1820 Gen2x1 Test Results

Test Equipment

RC: RK3588-EVB1-LP4X-V10 PCIe 3.0 4 lanes with main core at 2.4G and secondary core at 1.8G

EP: RK1820 PCIe 2.0 1 lane

Standard EP Version

RC - Rockchip_PCIE_EP_Standard_Card_20250523_V3.1.0.tar.gz

EP - RK RT-Thread SDK develop-v5.1.x branch

Test Parameters:

- Utilizing EP DMA engine
- task_num = 1

RK1820 DRAM Memory Test Results

DMA Size	4K
min/max	83us/139us

Notes:

- If the test is initiated by RC and uses the default EP DMA engine, the write test actually corresponds to EP DMA read, and the read test corresponds to EP DMA write.
- The test results are statistical values after 1 second of testing, excluding the bias influence of the initialization process.
- The test results represent the best case scenario under continuous packet transmission; please adjust the gaps and other loads according to requirements to construct the worst case scenario.

8.7 Compatibility Enumeration

Test Objective

Validate the compatibility and robustness of EP board detection, identification, and enumeration on different RCs.

Test Method

RK EP primarily interfaces with RCs:

EP Device	RC Device	Cold Boot	Wake from Suspend	Hot Boot	Confirmation Method
RK3588 EVB4 V11	RK3588 EVB1 V10	20/20	20/20	20/20	Linux: lspci grep 356a
RK3588 EVB4 V11	PC ubuntu	20/20	20/20	20/20	Linux: lspci grep 356a
RK3588 EVB4 V11	PC windows	20/20	20/20	20/20	Device Manager to confirm corresponding devices
RK3568 iotest device	RK3588 EVB1 V10	20/20	20/20	20/20	Linux: lspci grep 356a
RK3568 iotest device	PC ubuntu	20/20	20/20	20/20	Linux: lspci grep 356a

Notes:

- Actual products should cover as many target RC devices as possible.
- Cold Boot: PC power off, power on, press the POWER button.
- Wake from Suspend:
 - RK RC: Enter and exit suspend with the POWER button.
 - Ubuntu: Desktop versions can enter through the interface buttons, other scenarios to be confirmed.
 - Windows: CTRL+L key combination to enter, other scenarios to be confirmed.
- Hot Boot:
 - RK3588: Enter root privileges, input the reboot command in the shell interface.

- Ubuntu: Enter root privileges, input the reboot command in the shell interface.
- Windows: Under the desktop, ALT+F4 key combination to pop up reset options, select reset, some PCs support direct power button reset.

8.8 Compatibility EP Reset and RC Re-enumeration Detection

Test Objective

Test the robustness of the business under the scenario of "EP standalone reset", focusing on the following aspects:

- PCIe remove/rescan functionality
- Whether communication is restored normally
- Whether the business is thriving

Test Method

The following tests should be conducted under the condition that the above-mentioned functional tests are all normal, and compile the executable file of examples/pcie_auto_test, refer to the examples/pcie_auto_test/README.md document to deploy the test environment.

- PCIe remove/rescan functionality
 - Execute `./pcie_auto_test_rc 2 0` and `./pcie_auto_test_rc 3 0` on the RC side to select the 0th device for Remove and Rescan operations.
- Whether communication is restored normally
- Whether the business is thriving
 - Execute `./pcie_auto_test_rc 0 0` on the RC side to select the 0th device for API stability operations.

8.9 Functional OTA

Please refer to the "Single EP Card Universal RC" section to confirm whether it is a product form that requires quick startup, and confirm the storage solution used and perform corresponding OTA verification, mainly including:

- Single Storage
- Dual Storage
- EP Boot

For detailed reference, please see the "EP Card OTA" section.

9. Chips Connect Exception Handling

9.1 RC-side x86 Platform Initialization of PCIe Reports mmap Failed, Operation Not Permitted

```
root@rk:/home/rc-sample# ./pcie_test 10240
resource path = /sys/bus/pci/devices/0000:01:00.0/enable
resource path = /sys/bus/pci/devices/0000:01:00.0/resource0
mmap failed, ret = 0, Operation not permitted
ep dev num : 1
Segmentation fault (core dumped)
```

Answer: This issue arises because Ubuntu, starting from version 20.04, has introduced a Linux kernel lockdown feature. If Secure Boot is enabled in the BIOS, it activates the kernel lockdown, leading to mmap failure. To resolve this, one must enter the BIOS and disable Secure Boot to proceed normally.

9.2 HugePage Configuration

The RK platform supports Hugepage sizes of 2M, 32M, and 1GB, while other platforms support sizes based on the actual system. The configuration for a single Hugepage size must be larger than the size of a single send buffer and a single receive buffer in the SDK, and the number of Hugepages must exceed the required Buff quantity.

For example: 5 send Buffs, with each Buff being 4M, and 5 receive Buffs, with each Buff being 8M.

In this case, it is not possible to configure Hugepage as 2M; the minimum configuration required is 32M.

The number of Hugepages needed for send Buffs = $(\text{int}) (5 / (32\text{M} / 4\text{M})) + 1 = 1$

The number of Hugepages needed for receive Buffs = $(\text{int}) (5 / (32\text{M} / 8\text{M})) + 1 = 2$

The final configuration would be `default_hugepagesz=32M, hugepagesz=32M, hugepages=3`

9.2.1 X86 Platform

Check the size of Hugepagesize

```
rk@debian:~$ cat /proc/meminfo | grep Huge
HugePages_Total:      2
HugePages_Free:       2
HugePages_Rsvd:       0
HugePages_Surp:       0
Hugepagesize:        2048 kB
Hugetlb:             4096 kB
```

Modify the size to 1GB

- Modify the cmdline parameter, edit /etc/default/grub, add `default_hugepagesz=1G`

```
GRUB_CMDLINE_LINUX_DEFAULT="quiet default_hugepagesz=1G hugepagesz=1G hugepages=1"
GRUB_CMDLINE_LINUX=""
```

- Generate the configuration

```
/usr/sbin/grub-mkconfig -o /boot/grub/grub.cfg
```

- Restart the device `reboot` or `systemctl poweroff`
- Check the size and count, `cat /proc/meminfo | grep Hugepagesize`

```
rk@debian:~$ cat /proc/meminfo | grep Huge
HugePages_Total:      1
HugePages_Free:       1
HugePages_Rsvd:       0
HugePages_Surp:       0
Hugepagesize:        1048576 kB
Hugetlb:              1048576 kB
```

9.2.2 RK Platform

Kernel Enabling Hugepage Configuration

```
+++ b/arch/arm64/configs/rockchip_linux_defconfig
@@ -1,5 +1,10 @@
CONFIG_DEFAULT_HOSTNAME="localhost"
CONFIG_SYSVIPC=y
+CONFIG_HUGETLBFS=y
CONFIG_NO_HZ=y
CONFIG_HIGH_RES_TIMERS=y
CONFIG_PREEMPT_VOLUNTARY=y
```

Configuring the Default Size and Number of Hugepages:

Add three properties `default_hugepagesz=32M hugepagesz=32M hugepages=16` to the cmdline, configuring Hugepages to be 32M in size and the number available to be 16.

```
+++ b/arch/arm64/boot/dts/rockchip/rk3588-linux.dtsi
@@ -6,7 +6,7 @@

/ {
    chosen: chosen {
-        bootargs = "earlycon=uart8250,mmio32,0xfeb50000 console=ttyFIQ0
irqchip.gicv3_pseudo_nmi=0 clk_gate.always_on=1 pm_domains.always_on=1
root=PARTUUID=614e0000-0000 rw rootwait";
+        bootargs = "earlycon=uart8250,mmio32,0xfeb50000 console=ttyFIQ0
irqchip.gicv3_pseudo_nmi=0 clk_gate.always_on=1 pm_domains.always_on=1
root=PARTUUID=614e0000-0000 rw rootwait default_hugepagesz=32M hugepagesz=32M hugepages=16
";
    };

    cspmu: cspmu@fd10c000 {
```

Recompile the kernel and flash it. After booting, check `/proc/meminfo` to confirm if it has been added successfully.

```
[root@RK3588:/userdata]# cat /proc/meminfo | grep Huge
HugePages_Total:      16
HugePages_Free:       16
HugePages_Rsvd:        0
HugePages_Surp:        0
Hugepagesize:         32768 kB
Hugetlb:              1048576 kB
```

9.3 Encountering Kernel Errors or Bus Errors After Operating on the user cmd Area or ep info Area via API Interfaces

There are two main reasons that could lead to this issue:

1. Utilizing interfaces such as `memset` or `memcpy` to manipulate this block of memory.
2. The application's compilation optimization level exceeds `-O1`.

In the first scenario, review the code to remove interfaces like `memset` or `memcpy` and instead use a for loop for manipulation.

In the second scenario, you can reduce the optimization level or add the compilation parameter `-fno-tree-loop-optimize` to disable loop optimization.

9.4 Building the Windows Driver Compilation Environment

Current driver development and verification platform: **System Version Windows 10 2004, WDK Version 10.0.19041.685.**

1. Confirm System Version

Right-click on the computer properties to check the system version or press Win+R and enter winver to view, focusing mainly on the operating system internal version: 19042.1052.

Version	Windows 10 Professional
Version Number	20H2
Installation Date	2021/6/5
Operating System Build	19042.1052

2. Download and install VS2019 and WDK, refer to online tutorials for guidance.

3. Disable driver signature enforcement on the Windows host, here are three methods provided:

- Advanced Startup Options:

Restart the computer and enter the advanced startup options during the boot process (usually requires pressing F8 at startup).

In the advanced startup options menu, if you see "Disable driver signature enforcement" or a similar option, and this option is selected, then digital signature verification is turned off.

Using Command Prompt (only for temporary disablement):

- Open Command Prompt (as administrator).

Enter the following command:

```
bcdedit /set testsigning on
```

- Using Local Group Policy Editor (applicable for Windows Pro and higher versions):

Open "Local Group Policy Editor" (gpedit.msc).

Navigate to "Computer Configuration" -> "Administrative Templates" -> "System" -> "Driver Installation".

Review the setting for "Device: Allow installation of unsigned drivers". If this setting is enabled, it means that digital signature verification may be disabled.

Restart the computer, if you see the words "Test Mode" in the bottom right corner of the desktop, this indicates that digital signature verification is disabled. If digital signatures are not disabled, the driver will fail to install.

4. Load Project Compilation

The VS project file is located in the build/windows/rk_projects directory, simply open the rk_projects.sln.

The project rc_pcie is the driver project, supporting 4 configurations (different configurations are due to some APIs being deprecated after Windows 10 2004 version, requiring isolation through configurations):

- Debug_eq_win10_2004 : Debug environment, for systems greater than Windows 10 2004
- Debug_less_win10_2004 : Debug environment, for systems less than or equal to Windows 10 2004
- Release_eq_win10_2004 : Release environment, for systems greater than Windows 10 2004
- Release_less_win10_2004 : Release environment, for systems less than or equal to Windows 10 2004

The project examples is the test project, supporting Debug and Release configurations, used to verify if the driver is functioning properly.

After successful compilation, the driver and test files will be output in the build/windows/rk_projects/x64 directory.

9.5 Windows Application Execution Error: MSVCP140.dll and vcruntime140.dll Not Found

Download and install [Visual C++ Redistributable for Visual Studio 2015](#)

9.6 Error Occurs When Attempting to Modify Debug Program Settings in Windows Debugging Mode. The value is protected by Secure Boot policy and cannot be modified or deleted

Enter Windows 10 BIOS and disable Secure Boot by setting "secure boot" to "disable".

10. Chips Connect API Reference

For detailed reference, please consult the document
Rockchip_Developer_Guide_PCIE_EP_Standard_Card_APIs.